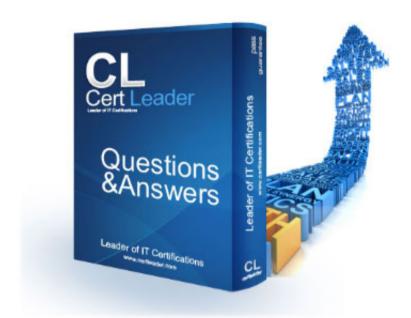# AZ-204 Dumps

Developing Solutions for Microsoft Azure (beta)

https://www.certleader.com/AZ-204-dumps.html

NEW QUESTION 1
- (Exam Topic 1)
You need to secure the Shipping Logic App. What should you use?

A. Azure App Service Environment (ASE)
B. Azure AD B2B integration
C. Integration Service Environment (ISE)
D. VNet service endpoint

Answer: C

Explanation:
Scenario: The Shipping Logic App requires secure resources to the corporate VNet and use dedicated storage resources with a fixed costing model.
You can access to Azure Virtual Network resources from Azure Logic Apps by using integration service environments (ISEs).
Sometimes, your logic apps and integration accounts need access to secured resources, such as virtual machines (VMs) and other systems or services, that are inside an Azure virtual network. To set up this access, you can create an integration service environment (ISE) where you can run your logic apps and create your integration accounts.
References:
https://docs.microsoft.com/en-us/azure/logic-apps/connect-virtual-network-vnet-isolated-environment-overview

NEW QUESTION 2
- (Exam Topic 1)
You need to migrate on-premises shipping data to Azure. What should you use?

A. Azure Migrate
B. Azure Cosmos DB Data Migration tool (dt.exe)
C. AzCopy
D. Azure Database Migration service

Answer: D

Explanation:
Migrate from on-premises or cloud implementations of MongoDB to Azure Cosmos DB with minimal downtime by using Azure Database Migration Service.
Perform resilient migrations of MongoDB data at scale and with high reliability.
Scenario: Data migration from on-premises to Azure must minimize costs and downtime.
The application uses MongoDB JSON document storage database for all container and transport information. References:
https://azure.microsoft.com/en-us/updates/mongodb-to-azure-cosmos-db-online-and-offline-migrations-are-now

NEW QUESTION 3
- (Exam Topic 1)
You need to support the message processing for the ocean transport workflow.
Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Step 1: Create an integration account in the Azure portal
You can define custom metadata for artifacts in integration accounts and get that metadata during runtime for your logic app to use. For example, you can provide metadata for artifacts, such as partners, agreements, schemas, and maps - all store metadata using key-value pairs.
Step 2: Link the Logic App to the integration account
A logic app that's linked to the integration account and artifact metadata you want to use. Step 3: Add partners, schemas, certificates, maps, and agreements
Step 4: Create a custom connector for the Logic App. References:
https://docs.microsoft.com/bs-latn-ba/azure/logic-apps/logic-apps-enterprise-integration-metadata
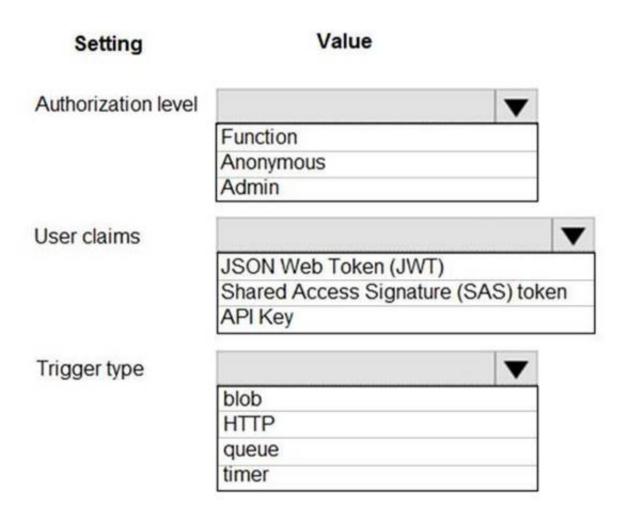
NEW QUESTION 4
- (Exam Topic 1)
You need to secure the Shipping Function app.
How should you configure the app? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

Setting	Value
Authorization level	▼
	Function
	Anonymous
	Admin
User claims	▼
	JSON Web Token (JWT)
	Shared Access Signature (SAS) token
	API Key
Trigger type	▼
	blob
	HTTP
	queue
	timer

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Scenario: Shipping Function app: Implement secure function endpoints by using app-level security and include Azure Active Directory (Azure AD).
Box 1: Function
Box 2: JSON based Token (JWT)
Azure AD uses JSON based tokens (JWTs) that contain claims Box 3: HTTP
How a web app delegates sign-in to Azure AD and obtains a token
User authentication happens via the browser. The OpenID protocol uses standard HTTP protocol messages. References:
https://docs.microsoft.com/en-us/azure/active-directory/develop/authentication-scenarios

NEW QUESTION 5
- (Exam Topic 3)
You develop a web application.
You need to register the application with an active Azure Active Directory (Azure AD) tenant.
Which three actions should you perform in sequence? To answer, move all actions from the list of actions to the answer area and arrange them in the correct order.

Actions

Select **Manifest** from the middle-tier service registration.

In Enterprise Applications, select **New application.**

Add a Cryptographic key.

Create a new application and provide the name, account type, and redirect URL

Select the Azure AD instance.

Use an access token to access the secure resource.

In App Registrations, select **New registration.**

Answer Area

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Register a new application using the Azure portal

≫ Sign in to the Azure portal using either a work or school account or a personal Microsoft account.

≫ If your account gives you access to more than one tenant, select your account in the upper right corner.
Set your portal session to the Azure AD tenant that you want.

≫ Search for and select Azure Active Directory. Under Manage, select App registrations.

≫ Select New registration. (Step 1)

≫ In Register an application, enter a meaningful application name to display to users.

≫ Specify who can use the application. Select the Azure AD instance. (Step 2)

≫ Under Redirect URI (optional), select the type of app you're building: Web or Public client (mobile & desktop). Then enter the redirect URI, or reply URL, for your application. (Step 3)

≫ When finished, select Register.

NEW QUESTION 6
- (Exam Topic 3)
You are developing an Azure Function App by using Visual Studio. The app will process orders input by an Azure Web App. The web app places the order information into Azure Queue Storage.
You need to review the Azure Function App code shown below. NOTE: Each correct selection is worth one point.

```
public static class orderProcessor
{
    [FunctionName("ProcessOrders")]
    public static void ProcessOrders([QueueTrigger("incoming-orders")]CloudQueueMessage myQueueItem, [Table
("Orders")][Collector<Order> tableBindings,
TraceWriter log)
    {
        log.Info($"Processing Order: {myQueueItem.Id}");
        log.Info($"Queue Insertion Time: {myQueueItem.InsertionTime}");
        log.Info($"Queue Expiration Time: {myQueueItem.ExpirationTime}");
        tableBindings.Add(JsonConvert.DeserializeObject<Order>(myQueueItem.AsString));
    }
    [FunctionName("ProcessOrders-Poison")]
    public static void ProcessFailedOrders([QueueTrigger("incoming-orders-poison")]CloudQueueMessage
myQueueItem, TraceWriter log)
    {
        logError($"Failed to process order: {myQueueItem.AsString}");
    }
}
```

	Yes	No
The code will log the time that the order was processed from the queue.	○	○
When the ProcessOrders function fails, the function will retry up to five times for a given order, including the first try.	○	○
When there are multiple orders in the queue, a batch of orders will be retrieved from the queue and the ProcessOrders function will run multiple instances concurrently to process the orders.	○	○
The ProcessOrders function will output the order to an Orders table in Azure Table Storage.	○	○

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: No
ExpirationTime - The time that the message expires. InsertionTime - The time that the message was added to the queue.
Box 2: Yes
maxDequeueCount - The number of times to try processing a message before moving it to the poison queue.
Default value is 5.
Box 3: Yes
When there are multiple queue messages waiting, the queue trigger retrieves a batch of messages and invokes function instances concurrently to process them.
By default, the batch size is 16. When the number being processed gets down to 8, the runtime gets another batch and starts processing those messages. So the maximum number of concurrent messages being processed per function on one virtual machine (VM) is 24.
Box 4: Yes References:
https://docs.microsoft.com/en-us/azure/azure-functions/functions-bindings-storage-queue

NEW QUESTION 7
- (Exam Topic 3)
You must ensure that the external party cannot access the data in the SSN column of the Person table.
Will each protection method meet the requirement? To answer, drag the appropriate responses to the correct protection methods. Each response may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Responses	Protection method	Response
Yes	Enable AlwaysOn encryption.	☐
No	Set the column encryption setting to disabled.	☐
	Assign users to the Public fixed database role.	☐
	Store column encryption keys in the system catalog view in the database.	☐

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: Yes
You can configure Always Encrypted for individual database columns containing your sensitive data. When setting up encryption for a column, you specify the information about the encryption algorithm and cryptographic keys used to protect the data in the column.
Box 2: No
Box 3: Yes
In SQL Database, the VIEW permissions are not granted by default to the public fixed database role. This enables certain existing, legacy tools (using older versions of DacFx) to work properly. Consequently, to work with encrypted columns (even if not decrypting them) a database administrator must explicitly grant the two VIEW permissions.
Box 4: No
All cryptographic keys are stored in an Azure Key Vault. References:
https://docs.microsoft.com/en-us/sql/relational-databases/security/encryption/always-encrypted-database-engine

NEW QUESTION 8
- (Exam Topic 3)

A company is developing a Java web app. The web app code is hosted in a GitHub repository located at https://github.com/Contoso/webapp.
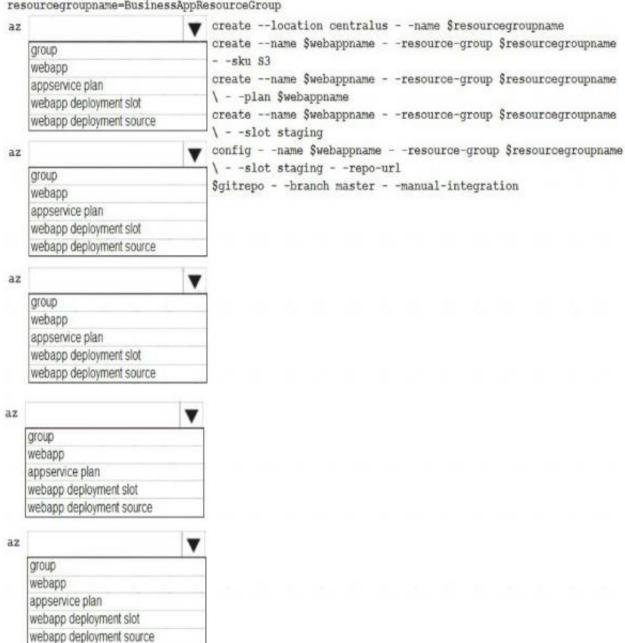The web app must be evaluated before it is moved to production. You must deploy the initial code release to a deployment slot named staging.
You need to create the web app and deploy the code.
How should you complete the commands? To answer, select the appropriate options in the answer area. NOTE: Each correct selection is worth one point.

```
gitrepo=https://github.com/Contoso/webapp
webappname=businesswebapp
resourcegroupname=BusinessAppResourceGroup
```

az [▼]	create --location centralus - -name $resourcegroupname
	create --name $webappname - -resource-group $resourcegroupname - -sku S3
group	create --name $webappname - -resource-group $resourcegroupname \ - -plan $webappname
webapp	create --name $webappname - -resource-group $resourcegroupname \ - -slot staging
appservice plan	
webapp deployment slot	
webapp deployment source	

az [▼]	config - -name $webappname - -resource-group $resourcegroupname \ - -slot staging - -repo-url
group	$gitrepo - -branch master - -manual-integration
webapp	
appservice plan	
webapp deployment slot	
webapp deployment source	

az [▼]
- group
- webapp
- appservice plan
- webapp deployment slot
- webapp deployment source

az [▼]
- group
- webapp
- appservice plan
- webapp deployment slot
- webapp deployment source

az [▼]
- group
- webapp
- appservice plan
- webapp deployment slot
- webapp deployment source

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: group
Create a resource group.
az group create --location westeurope --name myResourceGroup
Box 2: appservice plan
Create an App Service plan in STANDARD tier (minimum required by deployment slots). az appservice plan create --name $webappname --resource-group myResourceGroup --sku S1 Box 3: webapp
Create a web app.
az webapp create --name $webappname --resource-group myResourceGroup \
--plan $webappname
Box 4: webapp deployment slot
#Create a deployment slot with the name "staging".
az webapp deployment slot create --name $webappname --resource-group myResourceGroup \
--slot staging
Box 5: webapp deployment source
Deploy sample code to "staging" slot from GitHub.
az webapp deployment source config --name $webappname --resource-group myResourceGroup \
--slot staging --repo-url $gitrepo --branch master --manual-integration References:
https://docs.microsoft.com/en-us/azure/app-service/scripts/cli-deploy-staging-environment

NEW QUESTION 9
- (Exam Topic 3)
You are developing an application that use an Azure blob named data to store application data. The application creates blob snapshots to allow application state to be reverted to an earlier state. The Azure storage account has soft deleted enabled.
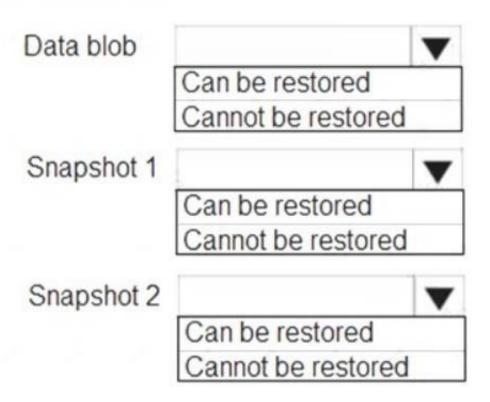The system performs the following operations in order:
•The blob is updated
•Snapshot 1 is created.
•Snapshot 2 is created.
•Snapshot 1 is deleted.
A system error then deletes the data blob and all snapshots. You need to determine which application states can be restored.
What is the restorability of the application data? To answer, select the appropriate options in the answer area. NOTE: Each correct selection is worth one point.

Application State   Restorability

Data blob

| Can be restored |
| Cannot be restored |

Snapshot 1

| Can be restored |
| Cannot be restored |

Snapshot 2

| Can be restored |
| Cannot be restored |

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: Can be restored
When enabled, soft delete enables you to save and recover your data when blobs or blob snapshots are deleted. This protection extends to blob data that is erased as the result of an overwrite.
Box 2: Cannot be restored It has been deleted.
Box 3: Can be restored It has not been deleted. References:
https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-soft-delete

NEW QUESTION 10
- (Exam Topic 3)
Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this question, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
Margie's Travel is an international travel and bookings management service. The company is expanding into restaurant bookings. You are tasked with implementing Azure Search for the restaurants listed in their solution.
You create the index in Azure Search.
You need to import the restaurant data into the Azure Search service by using the Azure Search .NET SDK. Solution:
* 1. Create a SearchIndexClient object to connect to the search index.
* 2. Create a DataContainer that contains the documents which must be added.
* 3. Create a DataSource instance and set its Container property to the DataContainer.
* 4. Call the Documents.Suggest method of the SearchIndexClient and pass the DataSource. Does the solution meet the goal?

A. Yes
B. No

Answer: B

Explanation:
Use the following method:
* 1. - Create a SearchIndexClient object to connect to the search index
* 2. - Create an IndexBatch that contains the documents which must be added.
* 3. - Call the Documents.Index method of the SearchIndexClient and pass the IndexBatch. References:
https://docs.microsoft.com/en-us/azure/search/search-howto-dotnet-sdk

NEW QUESTION 10
- (Exam Topic 3)
You are developing a project management service by using ASP.NET. The service hosts conversations, files, to-do lists, and a calendar that users can interact with at any time.
The application uses Azure Search for allowing users to search for keywords in the project data.
You need to implement code that creates the object which is used to create indexes in the Azure Search service.
Which two objects should you use? Each correct answer presents part of the solution. NOTE: Each correct selection is worth one point.

A. SearchService
B. SearchIndexClient
C. SearchServiceClient
D. SearchCredentials

Answer: BC

Explanation:
The various client libraries define classes like Index, Field, and Document, as well as operations like Indexes.Create and Documents.Search on the SearchServiceClient and SearchIndexClient classes.
Example:
The sample application we'll be exploring creates a new index named "hotels", populates it with a few documents, then executes some search queries. Here is the main program, showing the overall flow:
/ This sample shows how to delete, create, upload documents and query an index static void Main(string[] args)
{
IConfigurationBuilder builder = new ConfigurationBuilder().AddJsonFile("appsettings.json"); IConfigurationRoot configuration = builder.Build();
SearchServiceClient serviceClient = CreateSearchServiceClient(configuration); Console.WriteLine("{0}", "Deleting index...\n");
DeleteHotelsIndexIfExists(serviceClient);
Console.WriteLine("{0}", "Creating index...\n"); CreateHotelsIndex(serviceClient);
ISearchIndexClient indexClient = serviceClient.Indexes.GetClient("hotels"); References:
https://docs.microsoft.com/en-us/azure/search/search-howto-dotnet-sdk

NEW QUESTION 11
- (Exam Topic 3)
Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
Margie's Travel is an international travel and bookings management service. The company is expanding into restaurant bookings. You are tasked with implementing Azure Search tor the restaurants listed in their solution.
You create the index in Azure Search.
You need to import the restaurant data into the Azure Search service by using the Azure Search NET SDK. Solution:
* 1. Create a SearchServiceClient object to connect to the search index.
* 2. Create a DataContainer that contains the documents which must be added.
* 3. Create a DataSource instance and set its Container property to the DataContainer.
* 4. Set the DataSource property of the SearchServiceClient Does the solution meet the goal?

A. Yes
B. No

Answer: B

Explanation:
Use the following method:
* 1.Create a SearchIndexClient object to connect to the search index
* 2.Create an IndexBatch that contains the documents which must be added.
* 3. Call the Documents.Index method of the SearchIndexClient and pass the IndexBatch. References:
https://docs.microsoft.com/en-us/azure/search/search-howto-dotnet-sdk

NEW QUESTION 15
- (Exam Topic 3)
Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You are developing an Azure solution to collect point-of-sale (POS) device data from 2,000 stores located throughout the world. A single device can produce 2 megabytes (MB) of data every 24 hours. Each store location has one to five devices that send data.
You must store the device data in Azure Blob storage. Device data must be correlated based on a device identifier. Additional stores are expected to open in the future.
You need to implement a solution to receive the device data.
Solution: Provision an Azure Service Bus. Configure a topic to receive the device data by using a correlation filter.
Does the solution meet the goal?

A. Yes
B. No

Answer: A

Explanation:
A message is raw data produced by a service to be consumed or stored elsewhere. The Service Bus is for high-value enterprise messaging, and is used for order processing and financial transactions.
Reference:
https://docs.microsoft.com/en-us/azure/event-grid/compare-messaging-services

NEW QUESTION 18
- (Exam Topic 3)
You develop and deploy a Java RESTful API to Azure App Service.
You open a browser and navigate to the URL for the API. You receive the following error message:

```
Failed to load http://api.azurewebsites.net:6000/#/api/Products: No 'Access-
Control-Allow-Origin' header is present on the requested resource.
Origin 'http://localhost:6000' is therefore not allowed access
```

You need to resolve the error.
What should you do?

A. Bind an SSL certificate
B. Enable authentication

C. Enable CORS
D. Map a custom domain
E. Add a CDN

Answer: C

Explanation:
We need to enable Cross-Origin Resource Sharing (CORS). References:
https://medium.com/@xinganwang/a-practical-guide-to-cors-51e8fd329a1f

NEW QUESTION 23
- (Exam Topic 3)
Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You are developing an Azure solution to collect point-of-sale (POS) device data from 2,000 stores located throughout the world. A single device can produce 2 megabytes (MB) of data every 24 hours. Each store location has one to five devices that send data.
You must store the device data in Azure Blob storage. Device data must be correlated based on a device identifier. Additional stores are expected to open in the future.
You need to implement a solution to receive the device data.
Solution: Provision an Azure Event Grid. Configure event filtering to evaluate the device identifier. Does the solution meet the goal?

A. Yes
B. No

Answer: B

Explanation:
Instead use an Azure Service Bus, which is used order processing and financial transactions.
Note: An event is a lightweight notification of a condition or a state change. Event hubs is usually used reacting to status changes.
Reference:
https://docs.microsoft.com/en-us/azure/event-grid/compare-messaging-services

NEW QUESTION 28
- (Exam Topic 3)
You are implementing a software as a service (SaaS) ASP.NET Core web service that will run as an Azure Web App. The web service will use an on-premises SQL Server database for storage. The web service also includes a WebJob that processes data updates. Four customers will use the web service.

⟫ Each instance of the WebJob processes data for a single customer and must run as a singleton instance.

⟫ Each deployment must be tested by using deployment slots prior to serving production data.

⟫ Azure costs must be minimized.

⟫ Azure resources must be located in an isolated network. You need to configure the App Service plan for the Web App.
How should you configure the App Service plan? To answer, select the appropriate settings in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

App service plan setting	Value
Number of VM instances	▼
	2
	4
	8
	16
Pricing tier	▼
	Isolated
	Standard
	Premium
	Consumption

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Number of VM instances: 4
You are not charged extra for deployment slots. Pricing tier: Isolated
The App Service Environment (ASE) is a powerful feature offering of the Azure App Service that gives network isolation and improved scale capabilities. It is

essentially a deployment of the Azure App Service into a subnet of a customer's Azure Virtual Network (VNet).
References:
https://azure.microsoft.com/sv-se/blog/announcing-app-service-isolated-more-power-scale-and-ease-of-use/

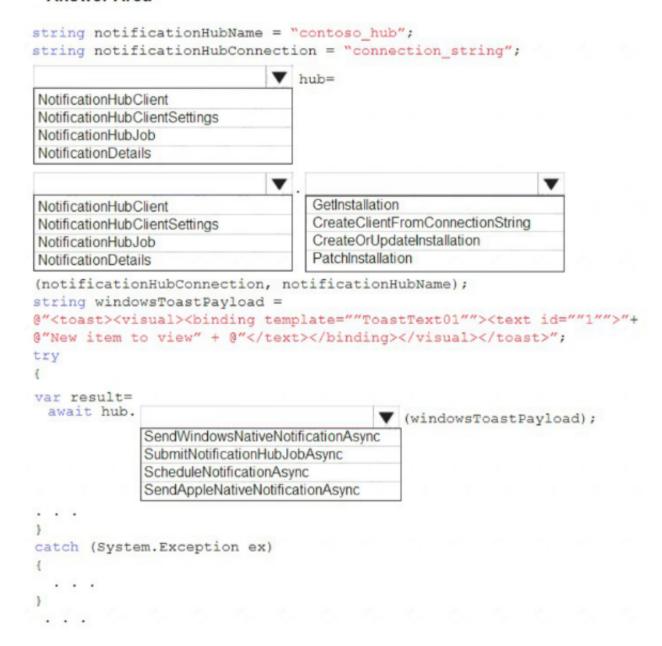**NEW QUESTION 32**
- (Exam Topic 3)
You develop a news and blog content delivery app for Windows devices.
A notification must arrive on a user's device when there is a new article available for them to view. You need to implement push notifications.
How should you complete the code segment? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

```
string notificationHubName = "contoso_hub";
string notificationHubConnection = "connection_string";
```

| ▼ | hub= |

```
NotificationHubClient
NotificationHubClientSettings
NotificationHubJob
NotificationDetails
```

| ▼ | . | | ▼ |

```
NotificationHubClient              GetInstallation
NotificationHubClientSettings      CreateClientFromConnectionString
NotificationHubJob                 CreateOrUpdateInstallation
NotificationDetails                PatchInstallation
```

```
(notificationHubConnection, notificationHubName);
string windowsToastPayload =
@"<toast><visual><binding template=""ToastText01""><text id=""1"">"+
@"New item to view" + @"</text></binding></visual></toast>";
try
{
    var result=
      await hub.
```

| ▼ | (windowsToastPayload); |

```
SendWindowsNativeNotificationAsync
SubmitNotificationHubJobAsync
ScheduleNotificationAsync
SendAppleNativeNotificationAsync
```

```
. . .
}
catch (System.Exception ex)
{
    . . .
}
. . .
```

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: NotificationHubClient
Box 2: NotificationHubClient
Box 3: CreateClientFromConnectionString
// Initialize the Notification Hub
NotificationHubClient hub = NotificationHubClient.CreateClientFromConnectionString(listenConnString, hubName);
Box 4: SendWindowsNativeNotificationAsync Send the push notification.
var result = await hub.SendWindowsNativeNotificationAsync(windowsToastPayload); References:
https://docs.microsoft.com/en-us/azure/notification-hubs/notification-hubs-push-notification-registration-manag
https://github.com/MicrosoftDocs/azure-docs/blob/master/articles/app-service-mobile/app-service-mobile-windo

NEW QUESTION 33
- (Exam Topic 3)
You are developing an ASP.NET Core Web API web service. The web service uses Azure Application Insights for all telemetry and dependency tracking. The web service reads and writes data to a database other than Microsoft SQL Server.
You need to ensure that dependency tracking works for calls to the third-party database.
Which two Dependency Telemetry properties should you store in the database? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

A. Telemetry.Context.Operation.Id
B. Tetemetry.Context.Cloud.RoleInstance
C. Telemetry.Id
D. Telemetry.ContextSession.Id
E. Telemetry.Name

Answer: AC

Explanation:
References:
https://docs.microsoft.com/en-us/azure/azure-monitor/app/custom-operations-tracking

NEW QUESTION 36
- (Exam Topic 3)
You are building a website to access project data related to terms within your organization. The website does not allow anonymous access. Authentication performed using an Azure Active Directory (Azure AD) app named internal.
The website has the following authentication requirements:
•Azure AD users must be able to login to the website.
•Personalization of the website must be based on membership in Active Directory groups. You need to configure the application's manifest to meet the authentication requirements.
How should you configure the manifest? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

```
{
    ...
    "appId": "d61126e3-089b-4adb-b721-
d5023213df7d",

    [▼] : "All",
    ┌──────────────────────┐
    │ "optionalClaims"     │
    │ "groupMembershipClaims" │
    └──────────────────────┘

    [▼] : true
    ┌──────────────────────┐
    │ "allowPublicClient"      │
    │ "oauth2Permissions"      │
    │ "requiredResourceAccess" │
    │ "oauth2AllowImplicitFlow" │
    └──────────────────────┘

    ...
}
```

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: groupMembershipClaims
Scenario: Personalization of the website must be based on membership in Active Directory groups. Group claims can also be configured in the Optional Claims section of the Application Manifest. Enable group membership claims by changing the groupMembershipClaim
The valid values are: "All" "SecurityGroup" "DistributionList" "DirectoryRole"
Box 2: oauth2Permissions
Scenario: Azure AD users must be able to login to the website.
oauth2Permissions specifies the collection of OAuth 2.0 permission scopes that the web API (resource) app exposes to client apps. These permission scopes may be granted to client apps during consent.

NEW QUESTION 37
- (Exam Topic 3)
You are a developer for a software as a service (SaaS) company that uses an Azure Function to process orders. The Azure Function currently runs on an Azure Function app that is triggered by an Azure Storage queue.
You are preparing to migrate the Azure Function to Kubernetes using Kubernetes-based Event Driven Autoscaling (KEDA).
You need to configure Kubernetes Custom Resource Definitions (CRD) for the Azure Function.
Which CRDs should you configure? To answer, drag the appropriate CRD types to the correct locations. Each CRD type may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Answer Area

CRD types	Setting	CRD type
Secret		
Deployment	Azure Function code	[ ]
ScaledObject	Polling interval	[ ]
TriggerAuthentication	Azure Storage connection string	[ ]

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: Deployment
To deploy Azure Functions to Kubernetes use the func kubernetes deploy command has several attributes that directly control how our app scales, once it is deployed to Kubernetes.
Box 2: ScaledObject
With --polling-interval, we can control the interval used by KEDA to check Azure Service Bus Queue for messages.
Example of ScaledObject with polling interval apiVersion: keda.k8s.io/v1alpha1
kind: ScaledObject metadata:
name: transformer-fn namespace: tt
labels:
deploymentName: transformer-fn spec:
scaleTargetRef: deploymentName: transformer-fn pollingInterval: 5
minReplicaCount: 0
maxReplicaCount: 100
Box 3: Secret
Store connection strings in Kubernetes Secrets. Example: to create the Secret in our demo Namespace:
create the k8s demo namespace kubectl create namespace tt
grab connection string from Azure Service Bus KEDA_SCALER_CONNECTION_STRING=$(az servicebus queue authorization-rule keys list \
-g $RG_NAME \
--namespace-name $SBN_NAME \
--queue-name inbound \
-n keda-scaler \
--query "primaryConnectionString" \
-o tsv)
create the kubernetes secret
kubectl create secret generic tt-keda-auth \
--from-literal KedaScaler=$KEDA_SCALER_CONNECTION_STRING \
--namespace tt Reference:
https://www.thinktecture.com/en/kubernetes/serverless-workloads-with-keda/

NEW QUESTION 42
- (Exam Topic 3)
You have an Azure Batch project that processes and converts files and stores the files in Azure storage. You are developing a function to start the batch job.
You add the following parameters to the function.

Parameter name	Description
fileTasks	a list of tasks to be run
jobId	the identifier that must be assigned to the job
outputContainerSasUrl	a storage SAS URL to store successfully converted files
failedContainerSasUrl	a storage SAS URL to store copies of files that failed to convert.

You must ensure that converted files are placed in the container referenced by the outputContainerSasUrl parameter. Files which fail to convert are places in the container referenced by the failedContainerSasUrl parameter.
You need to ensure the files are correctly processed.
How should you complete the code segment? To answer, select the appropriate options in the answer area. NOTE: Each correct selection is worth one point.

Answer Area

```
public list<CloudTasks> StartTasks(List<FileTask> fileTasks, string jobId,
  string outputContainerSasUrl, string failedContainerSasUrl)
{
  BatchSharedKeyCredentials sharedKeyCredentials =
    new BatchSharedKeyCredentials(batchAccountUrl, batchAccountName,
batchAccountKey);
  List<CloudTask> tasks = new List<CloudTask>();
  using (BatchClient batchClient = BatchClient.Open(sharedKeyCredentials))
  {
    CloudJob = batchClient.JobOperations. [▼] ();
```
```
                                          GetJob
                                          GetTask
                                          EnableJob
                                          CreateJob
```
```
        job.Id = jobId,
        job.PoolInfromation = new PoolInformation { PoolId = poolId };
        job.Commit();
        fileTasks.ForEach((fileTask) =>
        {
            string taskId = $"Task{DateTime.Now.ToFileTimeUtc().ToString()}";
            CloudTask task = new CloudTask (taskId, fileTask.Command);
            List<OutputFile> outputFileList = new List<OutputFile>();
            OutputFileBlobContainerDestination outputContainer =
              new OutputFileBlobContainerDestination(outputContainerSasUrl);
            OutputFileBlobContainerDestination failedContainer =
        new OutputFileBlobContainerDestination (failedContainerSasUrl);
      outputFileList.Add(new OutputFile(fileTask.Output,
          new OutputFileDestination(outputContainer),
          new OutputFileUploadOptions(OutputFileUploadCondition. [▼] )));
```
```
                                                              TaskSuccess
                                                              TaskFailure
                                                              TaskCompletion
```
```
        outputFileList.Add(new OutputFile(fileTask.Output,
          new OutputFileDestination(failedContainer),
          new OutputFileUploadOptions(OutputFileUploadCondition, [▼] )));
```
```
                                                              TaskSuccess
                                                              TaskFailure
                                                              TaskCompletion
```
```
        task [▼] =outputFileList;
```
```
            OutputFiles
            FilesToStage
            ResourceFiles
            StageFiles
```
```
        task.Add(task);
        });
        ]
        return tasks,
  }
}
```

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: CreateJob
Box 2: TaskSuccess
TaskSuccess: Upload the file(s) only after the task process exits with an exit code of 0.
Incorrect: TaskCompletion: Upload the file(s) after the task process exits, no matter what the exit code was. Box 3: TaskFailure
TaskFailure:Upload the file(s) only after the task process exits with a nonzero exit code. Box 4: OutputFiles
To specify output files for a task, create a collection of OutputFile objects and assign it to the CloudTask.OutputFiles property when you create the task.

References:
https://docs.microsoft.com/en-us/dotnet/api/microsoft.azure.batch.protocol.models.outputfileuploadcondition https://docs.microsoft.com/en-us/azure/batch/batch-task-output-files

NEW QUESTION 45
- (Exam Topic 3)
You are developing an application that uses Azure Storage Queues. You have the following code:

```
CloudStorageAccount storageAccount = CloudStorageAccount.Parse
(CloudConfigurationManager.GetSetting("StorageConnectionString"));
CloudQueueClient queueClient = storageAccount.CreateCloudQueueClient()

CloudQueue queue = queueClient.GetQueueReference("appqueue") ;
await queu.CreateIfNotExistsAsync() ;

CloudQueueMessage peekedMessage = await queue.PeekMessageAsync() ;
if (peekedMessage != null)
{
    Console.WriteLine("The peeked message is: {0}", peekedMessage.AsString);
}
CloudQueueMessage message = await queue.GetMessageAsync() ;
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

Answer Area

Statement	Yes	No
The code configures the lock duration for the queue.	○	○
The last message read remains in the queue after the code runs.	○	○
The storage queue remains in the storage account after the code runs.	○	○

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: No
The QueueDescription.LockDuration property gets or sets the duration of a peek lock; that is, the amount of time that the message is locked for other receivers.
The maximum value for LockDuration is 5 minutes; the default value is 1 minute.
Box 2: Yes
You can peek at the message in the front of a queue without removing it from the queue by calling the PeekMessage method.
Box 3: Yes Reference:
https://docs.microsoft.com/en-us/azure/storage/queues/storage-dotnet-how-to-use-queues https://docs.microsoft.com/en-us/dotnet/api/microsoft.servicebus.messaging.queuedescription.lockduration

NEW QUESTION 47
- (Exam Topic 3)
You have an app that stores player scores for an online game. The app stores data in Azure tables using a class named PlayerScore as the table entity. The table is populated with 100,000 records.
You are reviewing the following section of code that is intended to retrieve 20 records where the player score exceeds 15,000. (Line numbers are included for reference only.)

```
1  public void GetScore(string.playerId, int score, string gameName)
2  {
3      Table Query<DynamicTableEntity> query - new TableQuery<DynamicTableEntity>().Select(new.string[.] { "Score" })
           .Where(TableQuery.GenerateFilterConditionForInt("Score", QueryComparisons.GretaerThanOrEqual, 15000)).Take
(20);
4      EntityResolver<KeyValuePair<string, int?>> resolver =
        (partitionKey, rowKey, ts, props, etag) => new KeyValuePair<string, int?>(rowKey, props["Score"].Int32Value);
5      foreach (var scoreItem in scoreTable.ExecuteQuery(query, resolver, null, null))
6      {
        Console.Writeline($"{scoreItem.Key} {scoreItem.Value}");
7      }
8  }

9  public class PlayerScore : TableEntity
10 {
11     public PlayerScore(string gameId, string playerId, int score, long timePlayed)
12     {
13         PartitionKey = gameId;
14         RowKey = playerId;
15         Score = score;
16         TimePlayed = timePlayed;
17     }
18     public int Score { get; set; }
19     public long TimePlayed { get; set; }
20 }
```

You have the following code. (Line numbers are included for reference only.)
You store customer information in an Azure Cosmos database. The following data already exists in the database:

```
01  CloudTableClient tableClient = account.CreateCloudTableClient();
02  CloudTable table = tableClient.GetTableReference("people");
03  TableQuery<CustomerEntity> query = new TableQuery<CustomerEntity>()
04    .Where(TableQuery.CombineFilters(
05  TableQuery.Generate.And, TableQuery.GenerateFilterCondition(Email, QueryComparisons.Equal, "Smith")
06  TableOperstors.And, TableQuery.GenerateFilterCondition(Email, QueryComparisons.Equal,
"ssmith@contoso.com")
07  ));
08  await table.ExecuteQuerySegmentedAsync<CustomerEntity>(query, null);
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No. NOTE: Each correct selection is worth one point.

	Yes	No
The code queries the Azure table and retrieves the TimePlayed property from the table	O	O
The code will display a maximum of twenty records.	O	O
All records will be sent to the client. The client will display records for scores greater than or equal to 15,000.	O	O
The scoreItem.Key property of the KeyValuePairs that ExecuteQuery returns will contain a value for PlayerID.	O	O

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: No
Box 2: Yes
The TableQuery.Take method defines the upper bound for the number of entities the query returns. Example:
query.Take(10); Box 3: Yes
Box 4: Yes References:
https://www.vkinfotek.com/azureqa/how-do-i-query-azure-table-storage-using-tablequery-class.html

NEW QUESTION 48
- (Exam Topic 3)
You are preparing to deploy an ASP.NET Core website to an Azure Web App from a GitHub repository. The website includes static content generated by a script.
You plan to use the Azure Web App continuous deployment feature.
You need to run the static generation script before the website starts serving traffic.
What are two possible ways to achieve this goal? Each correct answer presents a complete solution. NOTE: Each correct selection is worth one point.

A. Create a file named .deployment in the root of the repository that calls a script which generates the static content and deploys the website.
B. Add a PreBuild target in the websites csproj project file that runs the static content generation script.
C. Create a file named run.cmd in the folder /run that calls a script which generates the static content and deploys the website.
D. Add the path to the static content generation tool to WEBSITE_RUN_FROM_PACKAGE setting in the host.json file.

Answer: AD

Explanation:
A: To customize your deployment, include a .deployment file in the repository root.
You just need to add a file to the root of your repository with the name .deployment and the content: [config]
command = YOUR COMMAND TO RUN FOR DEPLOYMENT
this command can be just running a script (batch file) that has all that is required for your deployment, like copying files from the repository to the web root directory for example.
D: In Azure, you can run your functions directly from a deployment package file in your function app. The other option is to deploy your files in the d:\home\site\wwwroot directory of your function app (see A above).
To enable your function app to run from a package, you just add a WEBSITE_RUN_FROM_PACKAGE setting to your function app settings.
Note: The host.json metadata file contains global configuration options that affect all functions for a function app.
References:
https://github.com/projectkudu/kudu/wiki/Custom-Deployment-Script
https://docs.microsoft.com/bs-latn-ba/azure/azure-functions/run-functions-from-deployment-package

NEW QUESTION 51
- (Exam Topic 3)
You are implementing a software as a service (SaaS) ASP.NET Core web service that will run as an Azure Web App. The web service will use an on-premises SQL Server database for storage. The web service also includes a WebJob that processes data updates. Four customers will use the web service.
•Each instance of the WebJob processes data for a single customer and must run as a singleton instance.
•Each deployment must be tested by using deployment slots prior to serving production data.
•Azure costs must be minimized.
•Azure resources must be located in an isolated network.
You need to configure the App Service plan for the Web App.
How should you configure the App Service plan? To answer, select the appropriate settings in the answer area. NOTE: Each correct selection is worth one point.

Answer Area

App service plan setting	Value
Number of VM instances	▼
	2
	4
	8
	16
Pricing tier	▼
	Isolated
	Standard
	Premium
	Consumption

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Number of VM instances: 4
You are not charged extra for deployment slots. Pricing tier: Isolated
The App Service Environment (ASE) is a powerful feature offering of the Azure App Service that gives network isolation and improved scale capabilities. It is essentially a deployment of the Azure App Service into a subnet of a customer's Azure Virtual Network (VNet).
References:
https://azure.microsoft.com/sv-se/blog/announcing-app-service-isolated-more-power-scale-and-ease-of-use/

NEW QUESTION 56
- (Exam Topic 3)
Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You develop a software as a service (SaaS) offering to manage photographs. Users upload photos to a web service which then stores the photos in Azure Storage Blob storage. The storage account type is
General-purpose V2.
When photos are uploaded, they must be processed to produce and save a mobile-friendly version of the image. The process to produce a mobile-friendly version of the image must start in less than one minute.
You need to design the process that starts the photo processing.
Solution: Convert the Azure Storage account to a BlockBlobStorage storage account. Does the solution meet the goal?

A. Yes

B. No

Answer: B

Explanation:
Not necessary to convert the account, instead move photo processing to an Azure Function triggered from the blob upload..
Azure Storage events allow applications to react to events. Common Blob storage event scenarios include image or video processing, search indexing, or any file-oriented workflow.
Note: Only storage accounts of kind StorageV2 (general purpose v2) and BlobStorage support event integration. Storage (general purpose v1) does not support integration with Event Grid.
Reference:
https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-event-overview

NEW QUESTION 61
- (Exam Topic 3)
You are using Azure Front Door Service.
You are expecting inbound files to be compressed by using Brotli compression. You discover that inbound XML files are not compressed. The files are 9 megabytes (MB) in size.
You need to determine the root cause for the issue.
To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

Statement	Yes	No
The file MIME type is supported by the service.	○	○
Edge nodes must be purged of all cache assets.	○	○
The compression type is supported.	○	○

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: No
Front Door can dynamically compress content on the edge, resulting in a smaller and faster response to your clients. All files are eligible for compression.
However, a file must be of a MIME type that is eligible for compression list.
Box 2: No
Sometimes you may wish to purge cached content from all edge nodes and force them all to retrieve new updated assets. This might be due to updates to your web application, or to quickly update assets that contain incorrect information.
Box 3: Yes
These profiles support the following compression encodings: Gzip (GNU zip), Brotli Reference:
https://docs.microsoft.com/en-us/azure/frontdoor/front-door-caching

NEW QUESTION 66
- (Exam Topic 3)
You are developing a web app that is protected by Azure Web Application Firewall (WAF). All traffic to the web app is routed through an Azure Application Gateway instance that is used by multiple web apps. The web app address is contoso.azurewebsites.net.
All traffic must be secured with SSL. The Azure Application Gateway instance is used by multiple web apps. You need to configure the Azure Application Gateway for the app.
Which two actions should you perform? Each correct answer presents part of the solution. NOTE: Each correct selection is worth one point.

A. In the Azure Application Gateway's HTTP setting, enable the Use for App service setting.
B. Convert the web app to run in an Azure App service environment (ASE).
C. Add an authentication certificate for contoso.azurewebsites.net to the Azure Application gateway.
D. In the Azure Application Gateway's HTTP setting, set the value of the Override backend path option to contoso22.azurewebsites.net.

Answer: AD

Explanation:
D: The ability to specify a host override is defined in the HTTP settings and can be applied to any back-end pool during rule creation.
The ability to derive the host name from the IP or FQDN of the back-end pool members. HTTP settings also provide an option to dynamically pick the host name from a back-end pool member's FQDN if configured with the option to derive host name from an individual back-end pool member.
A (not C): SSL termination and end to end SSL with multi-tenant services.
In case of end to end SSL, trusted Azure services such as Azure App service web apps do not require whitelisting the backends in the application gateway.

Therefore, there is no need to add any authentication certificates.

Add HTTP setting
saiappgw-appgw

* Protocol

| HTTP | **HTTPS** |

ℹ️ Authentication certificates are not required for trusted Azure certificates for end to end ssl to work

* Port ⓘ

| 443 | ✓ |

* Request timeout (seconds)

| 20 |

Override backend path ⓘ

| |

☑ Use for App service

☑ Use custom probe ⓘ

OK

Reference:
https://docs.microsoft.com/en-us/azure/application-gateway/application-gateway-web-app-overview

NEW QUESTION 67
- (Exam Topic 3)
You are creating an app that uses Event Grid to connect with other services. Your app's event data will be sent to a serverless function that checks compliance. This function is maintained by your company.
You write a new event subscription at the scope of your resource. The event must be invalidated after 3 specific period of time. You need to configure Event Grid to ensure security.
What should you implement? To answer, select the appropriate options in [he answer area. NOTE: Each correct selection is worth one point

Authentication   Type

WebHook event delivery

| ▼ |

| SAS tokens |
| Key authentication |
| JWT token |

Topic publishing

| ▼ |

| ValidationCode handshake |
| ValidationURL handshake |
| Management Access Control |

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: SAS tokens
Custom topics use either Shared Access Signature (SAS) or key authentication. Microsoft recommends SAS, but key authentication provides simple programming, and is compatible with many existing webhook publishers.
In this case we need the expiration time provided by SAS tokens. Box 2: ValidationCode handshake
Event Grid supports two ways of validating the subscription: ValidationCode handshake (programmatic) and ValidationURL handshake (manual).
If you control the source code for your endpoint, this method is recommended.

NEW QUESTION 68
- (Exam Topic 3)
You develop a gateway solution for a public facing news API. The news API back end is implemented as a RESTful service and uses an OpenAPI specification.
You need to ensure that you can access the news API by using an Azure API Management service instance. Which Azure PowerShell command should you run?

A. Import-AzureRmApiManagementApi –Context $ApiMgmtContext –SpecificationFormat "Swagger" -SpecificationPath $SwaggerPath –Path $Path
B. New-AzureRmApiManagementBackend -Context $ApiMgmtContext -Url $Url -Protocol http
C. New-AzureRmApiManagement –ResourceGroupName $ResourceGroup –Name $Name – Location $Location –Organization $Org –AdminEmail $AdminEmail
D. New-AzureRmApiManagementBackendProxy –Url $ApiUrl

Answer: D

Explanation:
New-AzureRmApiManagementBackendProxy creates a new Backend Proxy Object which can be piped when creating a new Backend entity.
Example: Create a Backend Proxy In-Memory Object
PS C:\>$secpassword = ConvertTo-SecureString "PlainTextPassword" -AsPlainText -Force
PS C:\>$proxyCreds = New-Object System.Management.Automation.PSCredential ("foo", $secpassword) PS C:\>$credential = New-AzureRmApiManagementBackendProxy -Url "http://12.168.1.1:8080"
-ProxyCredential $proxyCreds
PS C:\>$apimContext = New-AzureRmApiManagementContext -ResourceGroupName "Api-Default-WestUS" -ServiceName "contoso"
PS C:\>$backend = New-AzureRmApiManagementBackend -Context $apimContext -BackendId 123 -Url 'https://contoso.com/awesomeapi' -Protocol http -Title "first backend" -SkipCertificateChainValidation $true
-Proxy $credential -Description "backend with proxy server"
Creates a Backend Proxy Object and sets up Backend

NEW QUESTION 71
- (Exam Topic 3)
You are creating a CLI script that creates an Azure web app related services in Azure App Service. The web app uses the following variables:

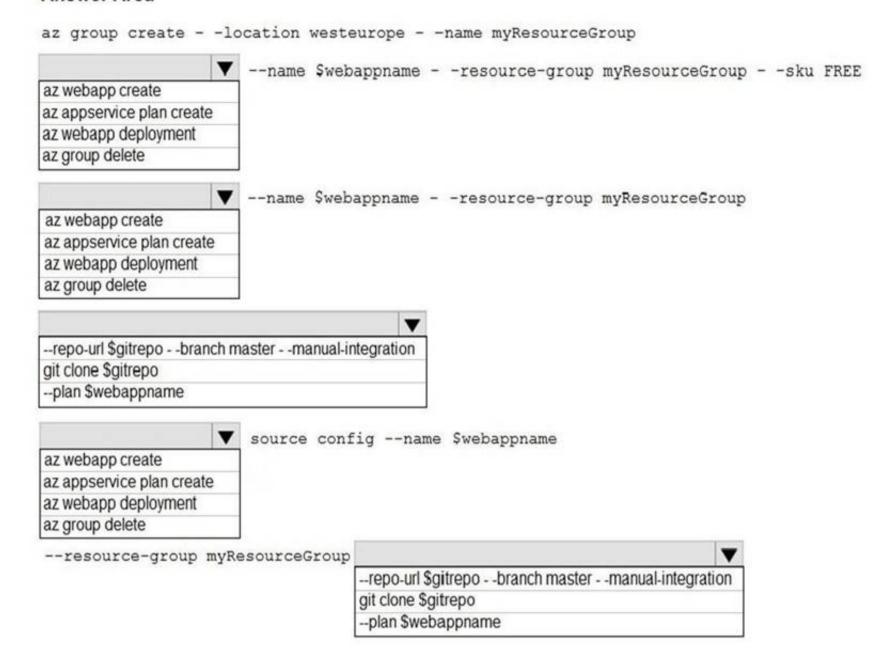| Variable name	Value
$gitrepo	https://github.com/Contos/webapp
$webappname	Webapp1103

You need to automatically deploy code from GitHub to the newly created web app.
How should you complete the script? To answer, select the appropriate options in the answer area. NOTE: Each correct selection is worth one point.

Answer Area

```
az group create - -location westeurope - -name myResourceGroup
```

▼	--name $webappname - -resource-group myResourceGroup - -sku FREE
az webapp create	
az appservice plan create	
az webapp deployment	
az group delete	

▼	--name $webappname - -resource-group myResourceGroup
az webapp create	
az appservice plan create	
az webapp deployment	
az group delete	

▼
--repo-url $gitrepo - -branch master - -manual-integration
git clone $gitrepo
--plan $webappname

▼	source config --name $webappname
az webapp create	
az appservice plan create	
az webapp deployment	
az group delete	

```
--resource-group myResourceGroup
```
▼
--repo-url $gitrepo - -branch master - -manual-integration
git clone $gitrepo
--plan $webappname

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: az appservice plan create
The azure group creates command successfully returns JSON result. Now we can use resource group to create a azure app service plan
Box 2: az webapp create Create a new web app..
Box 3: --plan $webappname
with the serviceplan we created in step 1. Box 4: az webapp deployment
Continuous Delivery with GitHub. Example:
az webapp deployment source config --name firstsamplewebsite1 --resource-group websites--repo-url $gitrepo
--branch master --git-token $token
Box 5: --repo-url $gitrepo --branch master --manual-integration Reference:
https://medium.com/@satish1v/devops-your-way-to-azure-web-apps-with-azure-cli-206ed4b3e9b1

NEW QUESTION 73
- (Exam Topic 3)
Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.
After you answer a question in this question, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You are developing a website that will run as an Azure Web App. Users will authenticate by using their Azure Active Directory (Azure AD) credentials.
You plan to assign users one of the following permission levels for the website: admin, normal, and reader. A user's Azure AD group membership must be used to determine the permission level. You need to configure authorization.
Solution:
•Create a new Azure AD application's manifest, set value of the groupMembershipClaims option to All.
•In the website, use the value of the groups claim from the JWI for the user to determine permissions. Does the solution meet the goal?

A. Yes
B. No

Answer: A

Explanation:
To configure Manifest to include Group Claims in Auth Token
* 1. Go to Azure Active Directory to configure the Manifest. Click on Azure Active Directory, and go to App registrations to find your application:
* 2. Click on your application (or search for it if you have a lot of apps) and edit the Manifest by clicking on it.
* 3. Locate the "groupMembershipClaims" setting. Set its value to either "SecurityGroup" or "All". To help you decide which:
"SecurityGroup" - groups claim will contain the identifiers of all security groups of which the user is a member.
"All" - groups claim will contain the identifiers of all security groups and all distribution lists of which the user is a member

Now your application will include group claims in your manifest and you can use this fact in your code. References:
https://blogs.msdn.microsoft.com/waws/2017/03/13/azure-app-service-authentication-aad-groups/

NEW QUESTION 75
- (Exam Topic 3)
You develop an ASP.NET Core MVC application. You configure the application to track webpages and custom events.
You need to identify trends in application usage.
Which Azure Application Insights Usage Analysis features should you use? To answer, drag the appropriate features to the correct requirements. Each feature may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Requirement	Feature
Which pages visited by users most often correlate to a product purchase?	
How does load time of the product display page affect a user's decision to purchase a product?	
Which events most influence a user's decision to continue to use the application?	
Are there places in the application that users often perform repetitive actions?	

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box1: Users Box 2: Impact
One way to think of Impact is as the ultimate tool for settling arguments with someone on your team about how slowness in some aspect of your site is affecting whether users stick around. While users may tolerate a certain amount of slowness, Impact gives you insight into how best to balance optimization and performance to maximize user conversion.
Box 3: Retention
The retention feature in Azure Application Insights helps you analyze how many users return to your app, and how often they perform particular tasks or achieve goals. For example, if you run a game site, you could compare the numbers of users who return to the site after losing a game with the number who return after winning. This knowledge can help you improve both your user experience and your business strategy.
Box 4: User flows
The User Flows tool visualizes how users navigate between the pages and features of your site. It's great for answering questions like:
How do users navigate away from a page on your site? What do users click on a page on your site?
Where are the places that users churn most from your site?
Are there places where users repeat the same action over and over?

NEW QUESTION 77
- (Exam Topic 3)
Your company has several websites that use a company logo image. You use Azure Content Delivery Network (CDN) to store the static image.
You need to determine the correct process of how the CDN and the Point of Presence (POP) server will distribute the image and list the items in the correct order.
In which order do the actions occur? To answer, move all actions from the list of actions to the answer area and arrange them in the correct order.

Actions

A user requests the image from the CDN URL. The DNS routes the request to the best performing POP location.

Subsequent requests for the file may be directed to the same POP using the CDN logo image URL. The POP edge server returns the files from cache if the TTL has not expired.

If no edge servers in the POP have the image in cache, the POP requests the file from the origin server.

The origin server returns the logo image to an edge server in the POP. An edge server in the POP caches the logo image and returns the image to the client.

Answer Area

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Step 1: A user requests the image..
A user requests a file (also called an asset) by using a URL with a special domain name, such as <endpoint name>.azureedge.net. This name can be an endpoint hostname or a custom domain. The DNS routes the request to the best performing POP location, which is usually the POP that is geographically closest to the user.
Step 2: If no edge servers in the POP have the..
If no edge servers in the POP have the file in their cache, the POP requests the file from the origin server. The origin server can be an Azure Web App, Azure Cloud Service, Azure Storage account, or any publicly accessible web server.
Step 3: The origin server returns the..
The origin server returns the file to an edge server in the POP.
An edge server in the POP caches the file and returns the file to the original requestor (Alice). The file remains cached on the edge server in the POP until the time-to-live (TTL) specified by its HTTP headers expires. If the origin server didn't specify a TTL, the default TTL is seven days.
Step 4: Subsequent requests for..
Additional users can then request the same file by using the same URL that the original user used, and can also be directed to the same POP.
If the TTL for the file hasn't expired, the POP edge server returns the file directly from the cache. This process results in a faster, more responsive user experience.
References:
https://docs.microsoft.com/en-us/azure/cdn/cdn-overview

NEW QUESTION 79
- (Exam Topic 3)
You are developing a back-end Azure App Service that scales based on the number of messages contained in a Service Bus queue.
A rule already exists to scale up the App Service when the average queue length of unprocessed and valid queue messages is greater than 1000.
You need to add a new rule that will continuously scale down the App Service as long as the scale up condition is not met.
How should you configure the Scale rule? To answer, select the appropriate options in the answer area. NOTE: Each correct selection is worth one point.

Answer Area

Scale rule ✕

Metric source

| Storage queue |
| Service Bus queue |
| Current resource |
| Storage queue (classic) |

Resource type

| Service Bus Namespaces ⌄ |

Resource

| MessageQueue1103 ⌄ |

* Queues

| itemqueue ⌄ |

Criteria

* Metric name

| Message Count |
| Active Message Count |

1 minute time grain

* Time grain statistic ❶

| Total |
| Maximum |
| Average |
| Count |

| Greater than |
| Greater than or equal to |
| Less than |
| Less than or equal to |

* Threshold

| 1000 |

Action

* Operation

| Increase count by |
| Increase count to |
| Decrease count by |
| Decrease count to |

* Instance count

| 1 |

* Cool down (minutes) ❶

| 5 |

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: Service bus queue
You are developing a back-end Azure App Service that scales based on the number of messages contained in a Service Bus queue.
Box 2: ActiveMessage Count
ActiveMessageCount: Messages in the queue or subscription that are in the active state and ready for delivery. Box 3: Count

Box 4: Less than or equal to
You need to add a new rule that will continuously scale down the App Service as long as the scale up condition is not met.
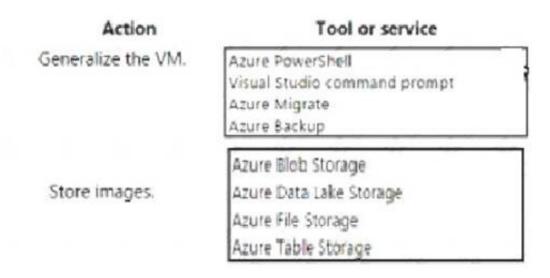Box 5: Decrease count by

NEW QUESTION 81
- (Exam Topic 3)
You are configuring a development environment for your team. You deploy the latest Visual Studio image from the Azure Marketplace to your Azure subscription. The development environment requires several software development kits (SDKs) and third-party components to support application development across the organization. You install and customize the deployed virtual machine (VM) for your development team. The customized VM must be saved to allow provisioning of a new team member development environment.
You need to save the customized VM for future provisioning.
Which tools or services should you use? To answer, select the appropriate options in the answer area. NOTE: Each correct selection is worth one point.

Answer Area

Action	Tool or service
Generalize the VM.	Azure PowerShell
	Visual Studio command prompt
	Azure Migrate
	Azure Backup
Store images.	Azure Blob Storage
	Azure Data Lake Storage
	Azure File Storage
	Azure Table Storage

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: Azure Powershell
Creating an image directly from the VM ensures that the image includes all of the disks associated with the VM, including the OS disk and any data disks.
Before you begin, make sure that you have the latest version of the Azure PowerShell module. You use Sysprep to generalize the virtual machine, then use Azure PowerShell to create the image. Box 2: Azure Blob Storage
References:
https://docs.microsoft.com/en-us/azure/virtual-machines/windows/capture-image-resource#create-an-image-of-a

NEW QUESTION 82
- (Exam Topic 3)
A company is developing a mobile app for field service employees using Azure App Service Mobile Apps as the backend.
The company's network connectivity varies throughout the day. The solution must support offline use and synchronize changes in the background when the app is online app.
You need to implement the solution.
How should you complete the code segment? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

```
var client = new MobileServiceClient("MOBILE_APP_URL");
var store = new MobileServiceSQLiteStore
(Constants.OfflineDbPath);
store.DefineTable<TodoItem>();
await client.SyncContext.IntitalizeAsync(store);
```

▼
var todoTable = client.GetSyncTable<TodoItem>();
var todoTable = client.GetTable<TodoItem>();
var todoTable = client.SyncTable;
var todoTable = client.Table;

```
await client.SyncContext.PushAsync();
```

▼
await todoTable.PullAsync("allTodoItems",todoTable.CreateQuery());
await todoTable.UpdateAsync();
todoTable.PullAsync("allTodoItems", todoTable.CreateQuery());
todoTable.UpdateAsync();

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: var todoTable = client GetSyncTable<TodoItem>()
To setup offline access, when connecting to your mobile service, use the method GetSyncTable instead of GetTable (example):
IMobileServiceSyncTable todoTable = App.MobileService.GetSyncTable(); / Box 2: await todoTable.PullAsync("allTodoItems",todo.Table.CreateQuery());
Your app should now use IMobileServiceSyncTable (instead of IMobileServiceTable) for CRUD operations. This will save changes to the local database and also keep a log of the changes. When the app is ready to synchronize its changes with the Mobile Service, use the methods PushAsync and PullAsync (example):
await App.MobileService.SyncContext.PushAsync(); await todoTable.PullAsync();
References:
https://azure.microsoft.com/es-es/blog/offline-sync-for-mobile-services/

NEW QUESTION 86
- (Exam Topic 3)
You develop a serverless application using several Azure Functions. These functions connect to data from within the code.
You want to configure tracing for an Azure Function App project. You need to change configuration settings in the hostjson file. Which tool should you use?

A. Azure portal
B. Azure PowerShell
C. Azure Functions Core Tools (Azure CLI)
D. Visual Studio

Answer: A

Explanation:
The function editor built into the Azure portal lets you update the function.json file and the code file for a function. The host.json file, which contains some runtime-specific configurations, is in the root folder of the function app.
References:
https://docs.microsoft.com/en-us/azure/azure-functions/functions-reference#fileupdate

NEW QUESTION 88
- (Exam Topic 3)
You are building a traffic monitoring system that monitors traffic along six highways. The system produces time series analysis-based reports for each highway.
Data from traffic sensors are stored in Azure Event Hub.
Traffic data is consumed by four departments. Each department has an Azure Web App that displays the time-series-based reports and contains a WebJob that processes the incoming data from Event Hub. All Web Apps run on App Service Plans with three instances.
Data throughout must be maximized. Latency must be minimized. You need to implement the Azure Event Hub.
Which settings should you use? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Setting	Value
Number of partitions	3 / 4 / 6 / 12 ▼
Partition Key	Highway / Department / Timestamp / VM name ▼

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: 6
The number of partitions is specified at creation and must be between 2 and 32. There are 6 highways.
Box 2: Highway References:
https://docs.microsoft.com/en-us/azure/event-hubs/event-hubs-features

NEW QUESTION 89
- (Exam Topic 3)
You are writing code to create and run an Azure Batch job. You have created a pool of compute nodes.
You need to choose the right class and its method to submit a batch job to the Batch service. Which method should you use?

A. JobOperations.CreateJobO
B. CloudJob.Enable(IEnumerable<BatchClientBehavior>)
C. CloudJob.CommitAsync(IEnumerable<BatchClientBehavior>, CancellationToken)
D. JobOperations.EnableJob(String, IEnumerable<BatchClientBehavior>)
E. JobOperations.EnableJobAsync(Strin
F. IEnumerable<BatchClientBehavior>. CancellationToken)

Answer: C

Explanation:
A Batch job is a logical grouping of one or more tasks. A job includes settings common to the tasks, such as priority and the pool to run tasks on. The app uses the BatchClient.JobOperations.CreateJob method to create a job on your pool.
The Commit method submits the job to the Batch service. Initially the job has no tasks.
{
CloudJob job = batchClient.JobOperations.CreateJob(); job.Id = JobId;
job.PoolInformation = new PoolInformation { PoolId = PoolId }; job.Commit();
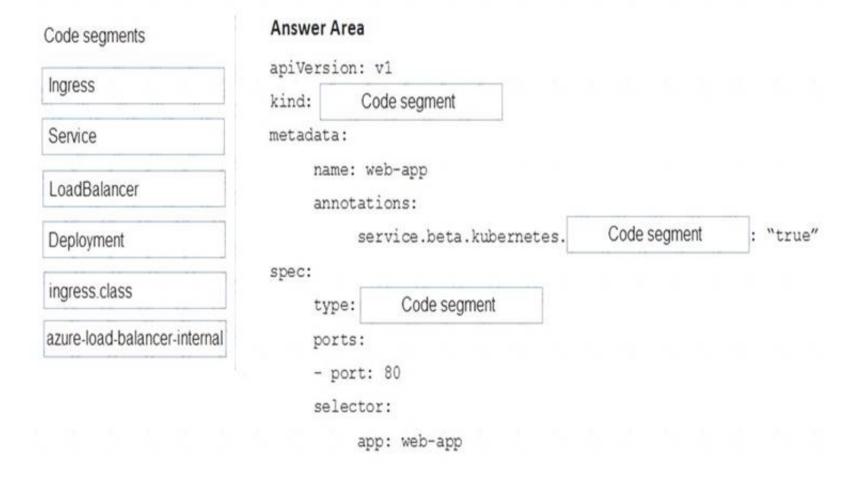}
References:
https://docs.microsoft.com/en-us/azure/batch/quick-run-dotnet

NEW QUESTION 93
- (Exam Topic 3)
You are preparing to deploy an application to an Azure Kubernetes Service (AKS) cluster. The application must only be available from within the VNet that includes the cluster. You need to deploy the application.
How should you complete the deployment YAML? To answer, drag the appropriate YAML segments to the correct locations. Each YAML segment may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
To create an internal load balancer, create a service manifest named internal-lb.yaml with the service type LoadBalancer and the azure-load-balancer-internal annotation as shown in the following example:
YAML:
apiVersion: v1 kind: Service metadata:
name: internal-app annotations:
service.beta.kubernetes.io/azure-load-balancer-internal: "true" spec:
type: LoadBalancer ports:
- port: 80 selector:
app: internal-app
References:
https://docs.microsoft.com/en-us/azure/aks/internal-lb

NEW QUESTION 97

- (Exam Topic 3)

A company uses Azure SQL Database to store data for an app. The data includes sensitive information.

You need to implement measures that allow only members of the managers group to see sensitive information. Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

A. Include the managers group.

B. Exclude the managers group.

C. Exclude the administrators group.

D. Navigate to the following URL:

```
PUT https://management.azure.com/subscriptions/00000000-1111-2222-3333-444444444444
/resourceGroups/rg01/providers/Microsoft.Sql/servers/server01/databases/customers
/transparentDataEncryption/current?api-version=2014-04-01
```

E. Run the following Azure PowerShell command:

```
New-AzureRmSqlDatabaseDataMaskingRule -SchemaName "dbo" -TableName "customers" `
-ColumnName "ssn" -MaskingFunction "Default"
```

A. Option A
B. Option B
C. Option C
D. Option D
E. Option E

Answer: BE

Explanation:

Dynamic data masking helps prevent unauthorized access to sensitive data by enabling customers to designate how much of the sensitive data to reveal with minimal impact on the application layer.

SQL users excluded from masking - A set of SQL users or AAD identities that get unmasked data in the SQL query results.

Note: The New-AzureRmSqlDatabaseDataMaskingRule cmdlet creates a data masking rule for an Azure SQL database.

References:

https://docs.microsoft.com/en-us/powershell/module/azurerm.sql/new-azurermsqldatabasedatamaskingrule?view

NEW QUESTION 102

- (Exam Topic 3)

You are developing a software solution for an autonomous transportation system. The solution uses large data sets and Azure Batch processing to simulate navigation sets for entire fleets of vehicles.

You need to create compute nodes for the solution on Azure Batch. What should you do?

A. In the Azure portal, create a Batch account.
B. In a .NET method, call the method: BatchClient.PoolOperations.CreatePool
C. In Python, implement the class: JobAddParameter
D. In Python, implement the class: TaskAddParameter

Answer: B

Explanation:

A Batch job is a logical grouping of one or more tasks. A job includes settings common to the tasks, such as priority and the pool to run tasks on. The app uses the BatchClient.JobOperations.CreateJob method to create a job on your pool.

NEW QUESTION 107

- (Exam Topic 3)

You are working for a company that designs mobile applications. They maintain a server where player records are assigned to their different games. The tracking system is new and in development.

The application uses Entity Framework to connect to an Azure Database. The database holds a Player table and Game table.

When adding a player, the code should insert a new player record, and add a relationship between an existing game record and the new player record.

The application will call CreatePlayerWithGame with the correct gameIdand the playerId to start the process. (Line numbers are included for reference only.)

```
01. namespace ContosoCradt
02. {
03    public class PlayerDbContext : DbContext
04.   {
05.   public PlayerDbContext() : base ("name-dBConnString") { ]
06.   public DbSet<Player> Players { get ; set ; }
07.   public DbSet<Game> Games { get ; set ]
08.   protected override void DmModelCreating{DBModelBuilder modelBuilder}
09.   {
10.    modelBuilder.Entity<Player>().MesMany(x => x.Games). WithMany (x => x Players);
11.   }
12.   }
13.    internal series class dbConfiguration : DbMigrationConfiguration<PlayerDbContext>
14.   {
15.   public dbConfiguration[] . {AutomaticMigrationsEnabled = true ; }
16.   {
17.    public class mp
18.    {
19.   public void CreatePlayerWithGame(int playerId, int gameId) => AddPlayer(playerId, GetGame[gameId]);
20.   public game GetGame(int gameId)
21.   {
22.    using (var db = new PlayerDbContext())
23.   {
24.    return db.Games.FirstOrDefault(x => x.GameId == gameId);
25.   }
26.    }
27.    public Player AddPlayer (int playerId, Game game)
28.    {
29.   using (var db = new PlayerDbContext())
30.   {
31.    var player = new Player
32.     {
33.       PlayerId = playerId,
34.       Games = new List <Game> {game },
35.   };
36.   db.Players.Add(player);
37.   db.SaveChanges();
38.   return player;
39.    }
40. }
41. public class Player
42. {
43. public int PlayerId { get ; set; }
44. public string PlayerName { get ; set; }
45. public virtual List<Game> Games { get ; set; }
46. }
47. public class Game
48. {
49. public int GameIs { get ; set }
50. public string Title { get ; set; }
51.  public string Platform { get ; set; ]
52.  public virtual List<Player> Players { get ; set; }
53.    }
54.  }
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No. NOTE: Each correct selection is worth one point.

	Yes	No
The code will successfully insert a player record.	O	O
The code has a bug and will insert an additional copy of the Game record with a new Id.	O	O
The code has a bug and will insert the wrong gameId value.	O	O
There is a valid many-to-many relationship between Players and Games.	O	O

A. Mastered

B. Not Mastered

Answer: A

Explanation:
Many-to-many relationships without an entity class to represent the join table are not yet supported. However, you can represent a many-to-many relationship by including an entity class for the join table and mapping two separate one-to-many relationships.
protected override void OnModelCreating(ModelBuilder modelBuilder)
{
modelBuilder.Entity<PostTag>() HasKey(t => new { t.PostId, t.TagId }); modelBuilder.Entity<PostTag>() HasOne(pt => pt.Post)
WithMany(p => p.PostTags) HasForeignKey(pt => pt.PostId); modelBuilder.Entity<PostTag>() HasOne(pt => pt.Tag) WithMany(t => t.PostTags) HasForeignKey(pt
=> pt.TagId);
}
}

NEW QUESTION 108
- (Exam Topic 3)
A company runs an international travel and bookings management service. The company plans to begin offering restaurant bookings. You must develop a solution that uses Azure Search and meets the following requirements:
• Users must be able to search for restaurants by name, description, location, and cuisine.
• Users must be able to narrow the results further by location, cuisine, rating, and family-friendliness.
• All words in descriptions must be included in searches. You need to add annotations to the restaurant class.
How should you complete the code segment? To answer, select the appropriate options in the answer area. NOTE: Each correct selection is worth one point.

```
[SerializePropertyNameAsCamelCase]
public class Restaurant
{
  [Key, IsFilterable]
  public int RestaurantId { get; set; }
  [IsSearchable, IsFilterable, IsSortable]
  public string Name { get; set; }
```
▼
| [IsSearchable.IsFilterable.IsSortable, IsFacetable] |
| [IsFilterable.IsFacetable,Required] |
| [IsSearchable] |
| [IsSearchable, Required] |

```
  public string location { get; set; }
  public string Phone { get; set; }
```
▼
| [Required] |
| [IsSearchable] |
| [IsFilterable, IsFacetable, Required] |
| [IsFilterable, IsFacetable, IsSortable] |

```
  public string Description { get; set; }
```
▼
| [IsFiltrable, IsSortable, IsSearchable] |
| [IsFilterable, IsSortable, IsFacetable] |
| [IsFiltrable, IsSortable, Key] |
| [IsFilteranle, IsSortable, IsSearchable, Required] |

```
  public double Rating { get; set; }
```
▼
| [IsSearchable, IsFilterable, IsFacetable] |
| [IsFilterbale, IsSortable, Key] |
| [IsFilterable, IsSortable, IsSearchable] |
| [IsFilterable, IsSortable, Key, Required] |

```
  public List<string> Cuisines { get; set; }
```
▼
| [IsFilterable, IsSortable, Key, Required] |
| [IsSearchable, IsSortable, IsFacetable] |
| [IsFilterable, IsSortable, Key, IsSearchable] |
| [IsFilterable, IsFacetable] |

```
  public bool FamilyFriendly { get; set; }
```

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: [IsSearchable.IsFilterable.IsSortable,IsFacetable] Location
Users must be able to search for restaurants by name, description, location, and cuisine.
Users must be able to narrow the results further by location, cuisine, rating, and family-friendliness. Box 2: [IsSearchable.IsFilterable.IsSortable,Required]
Description
Users must be able to search for restaurants by name, description, location, and cuisine. All words in descriptions must be included in searches.
Box 3: [IsFilterable,IsSortable,IsFaceTable] Rating

Users must be able to narrow the results further by location, cuisine, rating, and family-friendliness. Box 4: [IsSearchable.IsFilterable,IsFacetable]
Cuisines
Users must be able to search for restaurants by name, description, location, and cuisine.
Users must be able to narrow the results further by location, cuisine, rating, and family-friendliness. Box 5: [IsFilterable,IsFacetable]
FamilyFriendly
Users must be able to narrow the results further by location, cuisine, rating, and family-friendliness. References:
https://www.henkboelman.com/azure-search-the-basics/

NEW QUESTION 112
- (Exam Topic 3)
A company is developing a gaming platform. Users can join teams to play online and see leaderboards that include player statistics. The solution includes an entity named Team.
You plan to implement an Azure Redis Cache instance to improve the efficiency of data operations for entities that rarely change.
You need to invalidate the cache when team data is changed.
How should you complete the code? To answer, select the appropriate options in the answer area. NOTE: Each correct selection is worth one point.

```
void ClearCachedTeams()
{

    ▼
    IDatabase cache = Connection.GetDatabase();
    ICache cache = Connection.GetDatabase();


    ▼
    cache.KeyDelete("teams");
    cache.StringSet("teams",'');
    cache.ValueDelete("teams");
    cache.StringGet("teams", "");

    viewBag.nsg += Team data removed from cache. ";
}
```

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: IDatabase cache = connection.GetDatabase();
Connection refers to a previously configured ConnectionMultiplexer. Box 2: cache.StringSet("teams",")
To specify the expiration of an item in the cache, use the TimeSpan parameter of StringSet. cache.StringSet("key1", "value1", TimeSpan.FromMinutes(90));
References:
https://azure.microsoft.com/sv-se/blog/lap-around-azure-redis-cache-preview/

NEW QUESTION 116
- (Exam Topic 3)
You plan to create a Docker image that runs as ASP.NET Core application named ContosoApp. You have a setup script named setupScript.ps1 and a series of application files including ContosoApp.dll.
You need to create a Dockerfile document that meets the following requirements:
•Call setupScript.ps1 when the container is built.
•Run ContosoApp.dll when the container starts.
The Docker document must be created in the same folder where ContosoApp.dll and setupScript.ps1 are stored.
Which four commands should you use to develop the solution? To answer, move the appropriate commands from the list of commands to the answer area and arrange them in the correct order.

Commands

| RUN powershell ./setupScript.ps1
CMD ["dotnet", "ContosoApp.dll"] |

| EXPOSE ./ContosoApp/ /apps/ContosoApp |

| COPY ./. |

| FROM microsoft/aspnetcore:2.0 |

| WORKDIR /apps/ContosoApp |

| CMD powershell ./setupScript.ps1
ENTRYPOINT ["dotnet", "ContosoApp.dll"] |

Answer Area

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Step 1: WORKDIR /apps/ContosoApp Step 2: COPY ./The Docker document must be created in the same folder where ContosoApp.dll and setupScript.ps1 are stored.
Step 3: EXPOSE ./ContosApp/ /app/ContosoApp Step 4: CMD powershell ./setupScript.ps1 ENTRYPOINT ["dotnet", "ContosoApp.dll"]
You need to create a Dockerfile document that meets the following requirements:

> Call setupScript.ps1 when the container is built.

> Run ContosoApp.dll when the container starts. References:
https://docs.microsoft.com/en-us/azure/app-service/containers/tutorial-custom-docker-image

NEW QUESTION 118
- (Exam Topic 3)
You are developing an application. You have an Azure user account that has access to two subscriptions. You need to retrieve a storage account key secret from Azure Key Vault.
In which order should you arrange the PowerShell commands to develop the solution? To answer, move all commands from the list of commands to the answer area and arrange them in the correct order.

Powershell commands

| $secretvalue = ConvertTo-SecureString
$storAcctkey -AsPlainText
-Force
    Set-AzKeyVaultSecret -VaultName
$vaultName -Name $secretName
-SecretValue $secretvalue |

| Get-AzStorageAccountKey -
ResourceGroupName $resGroup -Name
$storAcct |

| Set-AzContext -SubscriptionId
$subscriptionID |

| Get-AzKeyVaultSecret -VaultName
$vaultName |

| Get-AzSubscription |

Answer Area

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Step 1: Get-AzSubscription
If you have multiple subscriptions, you might have to specify the one that was used to create your key vault. Enter the following to see the subscriptions for your account:
Get-AzSubscription
Step 2: Set-AzContext -SubscriptionId
To specify the subscription that's associated with the key vault you'll be logging, enter: Set-AzContext -SubscriptionId <subscriptionID>
Step 3: Get-AzStorageAccountKey You must get that storage account key.
Step 4: $secretvalue = ConvertTo-SecureString <storageAccountKey> -AsPlainText -Force
Set-AzKeyVaultSecret -VaultName <vaultName> -Name <secretName> -SecretValue $secretvalue After retrieving your secret (in this case, your storage account key), you must convert that key to a secure
string, and then create a secret with that value in your key vault.
Step 5: Get-AzKeyVaultSecret
Next, get the URI for the secret you created. You'll need this URI in a later step to call the key vault and retrieve your secret. Run the following PowerShell command and make note of the ID value, which is the secret's URI:
Get-AzKeyVaultSecret –VaultName <vaultName> Reference:
https://docs.microsoft.com/bs-latn-ba/Azure/key-vault/key-vault-key-rotation-log-monitoring
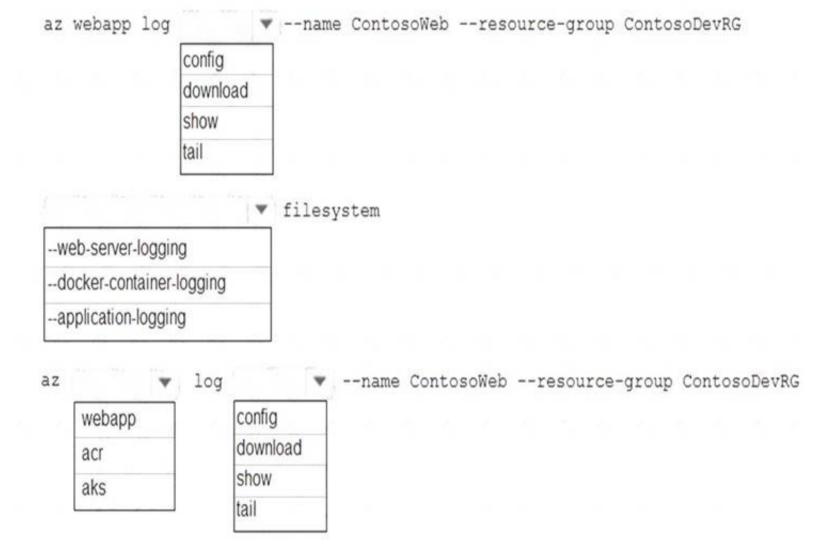
NEW QUESTION 121
- (Exam Topic 3)
You plan to deploy a web app to App Service on Linux. You create an App Service plan. You create and push a custom Docker image that image that contains the web app to Azure Container Registry.
You need to access the console logs generated from inside the container in real-time.
How should you complete the Azure CLI command? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: config
To Configure logging for a web app use the command: az webapp log config
Box 2: --docker-container-logging Syntax include:
az webapp log config [--docker-container-logging {filesystem, off}] Box 3: webapp
To download a web app's log history as a zip file use the command: az webapp log download
Box 4: download References:
https://docs.microsoft.com/en-us/cli/azure/webapp/log

NEW QUESTION 124
- (Exam Topic 3)
Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution. Determine whether the

solution meets the stated goals.
You are developing and deploying several ASP.Net web applications to Azure App Service. You plan to save session state information and HTML output. You must use a storage mechanism with the following requirements:
• Share session state across all ASP.NET web applications
• Support controlled, concurrent access to the same session state data for multiple readers and a single writer
• Save full HTTP responses for concurrent requests You need to store the information.
Proposed Solution: Deploy and configure Azure Cache for Redis. Update the web applications.
.
Does the solution meet the goal?

A. Yes
B. No

Answer: B

NEW QUESTION 125
- (Exam Topic 3)
ASP.NET Core API app by using C#. The API app will allow users to authenticate by using Twitter and Azure Active Directory (Azure AD).
Users must be authenticated before calling API methods. You must log the user's name for each method call. You need to configure the API method calls. Which values should you use? To answer, select the appropriate options in the answer area. NOTE: Each correct selection is worth one point.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Box 1: Authorize
Box 2: X-MS-CLIENT-PRINCIPAL-NAME
App Service passes user claims to your application by using special headers. External requests aren't allowed to set these headers, so they are present only if set by App Service. Some example headers include:
X-MS-CLIENT-PRINCIPAL-NAME X-MS-CLIENT-PRINCIPAL-ID
Here's the set of headers you get from Easy Auth for a Twitter authenticated user:
{
"cookie": "AppServiceAuthSession=Lx43...xHDTA==", "x-ms-client-principal-name": "evilSnobu",
"x-ms-client-principal-id": "35....", "x-ms-client-principal-idp": "twitter",
"x-ms-token-twitter-access-token": "35...Dj",
"x-ms-token-twitter-access-token-secret": "OK3...Jx",
}
References:
https://docs.microsoft.com/en-us/azure/app-service/app-service-authentication-how-to

NEW QUESTION 129
- (Exam Topic 3)
You develop Azure solutions.
You must connect to a No-SQL globally-distributed database by using the .NET API. You need to create an object to configure and execute requests in the database. Which code segment should you use?

A. new Container(EndpointUri, PrimaryKey);
B. new Database(Endpoint, PrimaryKey);
C. new CosmosClient(EndpointUri, PrimaryKey);

Answer: C

Explanation:
Example:

```
// Create a new instance of the Cosmos Client
this.cosmosClient = new CosmosClient(EndpointUri, PrimaryKey)
//ADD THIS PART TO YOUR CODE
await this.CreateDatabaseAsync(); Reference:
```
https://docs.microsoft.com/en-us/azure/cosmos-db/sql-api-get-started

NEW QUESTION 132
- (Exam Topic 3)
You are creating a hazard notification system that has a single signaling server which triggers audio and visual alarms to start and stop.
You implement Azure Service Bus to publish alarms. Each alarm controller uses Azure Service Bus to receive alarm signals as part of a transaction. Alarm events must be recorded for audit purposes. Each transaction record must include information about the alarm type that was activated.
You need to implement a reply trail auditing solution.
Which two actions should you perform? Each correct answer presents part of the solution. NOTE: Each correct selection is worth one point.

A. Assign the value of the hazard message SessionID property to the SequenceNumber property.
B. Assign the value of the hazard message SequenceNumber property to the DeliveryCount propert
C. Assign the value of the hazard message MessageId property to the DeliveryCount property.
D. Assign the value of the hazard message SessionID property to the ReplyToSessionId property.
E. Assign the value of the hazard message MessageId property to the SequenceNumber property.
F. Assign the value of the hazard message MessageId property to the CorrelationId property.

Answer: AB

NEW QUESTION 134
......

Thank You for Trying Our Product

* 100% Pass or Money Back

  All our products come with a 90-day Money Back Guarantee.

* One year free update

  You can enjoy free update one year. 24x7 online support.

* Trusted by Millions

  We currently serve more than 30,000,000 customers.

* Shop Securely

  All transactions are protected by VeriSign!

100% Pass Your AZ-204 Exam with Our Prep Materials Via below:

https://www.certleader.com/AZ-204-dumps.html